Przemysław Uznański 28/10/2019

Algorithms for Big Data

Fall Semester 2019

Exercise Set 4

Exercise 1:
Show that Cauchy distribution $f(x) = \frac{1}{\pi(1+x^2)}$ is 1-stable, that is for $X, Y \sim \mathsf{Cauchy}$ and $a, b \in \mathbb{R}$ we have $a * X + b * Y \sim (|a| + |b|) \cdot \mathsf{Cauchy}$.

Exercise 2:
Modify the algorithm for computing F_p of a stream (the one working for any p in the sequence of values) to work in semi-turnstile streams (updates (x_i, c_i) where $c_i \in \mathbb{R}^+$).

The goal of next few exercises is to show that p-stable distribution approach can be used for sketching of Hamming norm/distance. We follow the analysis from *"Comparing Data Streams Using Hamming Norms (How to Zero In)"* by Cormode, Datar, Indyk and Muthukrishnan (VLDB'02).[1]

Exercise 3:
Assume we operate in universe U of magnitude u, that is we have a promise that all our values we are ever going to see when sketching are integers from $\{-u, \ldots, 1, 0, 1, \ldots, u\}$. Show that for sufficiently small value p, F_p is an $1 \pm \varepsilon$ approximation of F_0. Show that $p = \Theta(\varepsilon / \log u)$ is small enough.
Show that u-factor approximation of L_p approximation is enough to obtain $1 \pm \varepsilon$ approximation of F_p, and thus $1 \pm 2\varepsilon$ of F_0.

Useful fact: when $p \to 0$, we have $\Pr(|X| > t) = \Theta(t^{-p})$ for X drawn from p-stable, 0 mean, normalized by median stable distribution.

Exercise 4: (2 pts)
Take value of p from Exercise 3 and desired level of approximation. How many samples do we need to take from this p-stable distribution to reach this level of approximation for median estimation? What is the memory complexity of our algorithm? Take into account the actual bit-size of bignums needed in our algorithm.[2]

[1] With the only change that we actualy want to do the correct analysis - the analysis in the paper has an error.

[2] And here is the omission of the authors...