

Algorithms for Big Data

Fall Semester 2019

Exercise Set 4

Exercise 1:

Show that Cauchy distribution $f(x) = \frac{1}{\pi(1+x^2)}$ is 1-stable, that is for $X, Y \sim \text{Cauchy}$ and $a, b \in \mathbb{R}$ we have $a * X + b * Y \sim (|a| + |b|) \cdot \text{Cauchy}$.

Exercise 2:

Modify the algorithm for computing F_p of a stream (the one working for any p in the sequence of values) to work in semi-turnstile streams (updates (x_i, c_i) where $c_i \in \mathbb{R}^+$).

The goal of next few exercises is to show that p -stable distribution approach can be used for sketching of Hamming norm/distance. We follow the analysis from "*Comparing Data Streams Using Hamming Norms (How to Zero In)*" by Cormode, Datar, Indyk and Muthukrishnan (VLDB'02).¹

Exercise 3:

Assume we operate in universe U of magnitude u , that is we have a promise that all our values we are ever going to see when sketching are integers from $\{-u, \dots, 1, 0, 1, \dots, u\}$. Show that for sufficiently small value p , F_p is an $1 \pm \varepsilon$ approximation of F_0 . Show that $p = \Theta(\varepsilon / \log u)$ is small enough.

Show that u -factor approximation of L_p approximation is enough to obtain $1 \pm \varepsilon$ approximation of F_p , and thus $1 \pm 2\varepsilon$ of F_0 .

Useful fact: when $p \rightarrow 0$, we have $\Pr(|X| > t) = \Theta(t^{-p})$ for X drawn from p -stable, 0 mean, normalized by median stable distribution.

Exercise 4:

(2 pts)

Take value of p from Exercise 3 and desired level of approximation. How many samples do we need to take from this p -stable distribution to reach this level of approximation for median estimation? What is the memory complexity of our algorithm? Take into account the actual bit-size of bignums needed in our algorithm.²

¹With the only change that we actually want to do the correct analysis - the analysis in the paper has an error.

²And here is the omission of the authors...