University of Wrocław: Algorithms for Big Data (Fall'19) 28/10/2019

Lecture 4: Point queries, heavy hitters

Lecturer: *Przemysław Uznański* Scribe: *Michał Jaworski*

1 Finding frequent elements

1.1 Majority elements (Boyer Moore 1981)

Definition 1. *An element of a multiset X, $|X| = n$ is called a* majority element *if it occurs more than $\frac{n}{2}$ times.*

We derive algorithm for finding majority element in 2-passes over the input. An useful fact:

Lemma 2. *If t is the majority element of X, and for some $x, y \in X$ we have $x \neq y$, then t is the majority element of $X \setminus \{x, y\}$.*

We can derive an algorithm utilizing this fact (this describes first pass, which returns a candidate for majority element).

```
s = anything
c = 0
for x in X:
    if x == s:
        c++
    else if c == 0:
        x = s
        c = 1
    else:
        c--
return s
```

If X contains a majority element it will be returned, otherwise we get random garbage. Second pass is just to verify if candidate is actually a majority element.

1.2 $\frac{1}{k}$-heavy elements

Definition 3. *An element of a multiset X, $|X| = n$, is called $\frac{1}{k}$-heavy if it occurs at least $\frac{n}{k}$ times.*

Lemma 4. *If t is an $\frac{1}{k}$-heavy element of X,and $x_1, ..., x_k \in X$ are pairwise distinct, then t is a $\frac{1}{k}$-heavy element of $X \setminus \{x_1, ..., x_k\}$.*

We can again derive an algorithm:

- Summary of a stream: a multiset S (represented for example as a collection of pairs (x_i, c_i))
- Invariant: $\#distinct(S) \leq k$

- Insert: add the element to S, then run the pruning step
- Pruning: while $\#distinct(S) > k$: remove from S one copy of each element in S

At the end the summary will contain all heavy hitters but it may also contain other elements - we need a second pass to check. This algorithm is deterministic and the summaries are mergeable.

2 L_p point queries, L_p heavy hitters

Assume turnstile streaming model, we maintain vector x under updates. Choose a norm L_p.

- **Point query**: for a given i return $x_i \pm \varepsilon |x|_p$
- **Heavy hitters**: $HH_\varepsilon(x) = \{i : |x_i| \geq \varepsilon |x|_p\}$. Output L such that:
  1. $HH_\varepsilon(x) \subseteq L \subseteq HH_{\varepsilon'}(x)$ for some $\varepsilon' < \varepsilon$, or
  2. $HH_\varepsilon(x) \subseteq L$ and $|L| = O(\frac{1}{\varepsilon^p})$ (intuition: $|HH_\varepsilon(x)| \leq \frac{1}{\varepsilon^p}$).

Observation 5. ***Roughly****, if we can solve one of these problems, then we can also solve the other (they reduce to each other).*

3 CountMin (Cormode, Muthukrishnan 2004)

3.1 Point queries

Assume $\forall_i x_i \geq 0$. Let $h : [n] \to [t]$ be a 2-wise independent hashing function for $t = \frac{2}{\varepsilon}$. Maintain an array $Z[1..t]$ such that

$$Z[j] = \sum_{i:h(i)=j} x_i$$

- Update(i, Δ): $Z[h(i)] \leftarrow Z[h(i)] + \Delta$
- Query(i): output $x'_i = Z[h(i)]$

Properties of queries:

1. $x'_i \geq x_i$
2. $\mathbb{E}[x'_i - x_i] = \sum_{j \neq i} \Pr[h(j) = h(i)] \cdot x_j \leq \frac{1}{t}|x|_1 = \frac{\varepsilon}{2}|x|_1$

From Markov's inequality:

$$\Pr[x'_i - x_i \geq \varepsilon |x|_1] \leq \frac{1}{2}$$

so we get an L_1 guarantee.

We can now derive the actual algorithm: we repeat the process independently $r = \log(\delta^{-1})$ times, using independent hashing functions $h_1, ..., h_r : [n] \to [t]$. We take the minimum of query results as the answer:

- Update(i, Δ): $\forall_{j \in [r]} Z[j][h_j(i)] \leftarrow Z[j][h_j(i)] + \Delta$
- Query(i): output $x'_i = \min_j Z[j][h_j(i)]$

We now get:

$$P(x_i' - x_i \geq \varepsilon |x|_1) \leq \left(\frac{1}{2}\right)^r = \delta$$

Recall we assumed $\forall_i x_i \geq 0$. For the general case replace minimum with median, and take $t = \frac{3}{\varepsilon}$. Finally, space complexity: $O\left(\frac{\log \delta^{-1}}{\varepsilon}\right)$ words and time complexity: $O(\log \delta^{-1})$ per update/query.

3.2 Heavy hitters

3.2.1 Generic transformation

This method only works in the semi-turnstile model.

- Maintain $|x|$ from sketch, keep heavy hitters in a priority queue
- On update run a point query, if x_i is a heavy hitter - insert it into the queue or update its weight
- Whenever the element with lowest weight gets below the $\varepsilon|x|$ threshold, remove it from the queue

3.2.2 Our case

We construct a binary search tree with $O(\log n)$ levels and n nodes in the lowest level. Each node represents the sum of values from an interval of the form $[a2^b + 1, (a+1)2^b]$. Each level is implemented as a separate CountMin structure. We perform queries recursively descending only into nodes where the output is at least $\varepsilon|x|_1$. Since there are only $\frac{1}{\varepsilon}$ such nodes at each level, we can apply the union bound $\delta' = \frac{\delta}{2\log n}$.

4 CountSketch (Charikar, Chen, Farach-Colton 2002)

Let $h_1, ..., h_r : [n] \to [t]$ and $s_1, ..., s_r : [n] \to \{-1, 1\}$ be 2-wise independent hashing functions.

- Update(i, Δ): $\forall_{j \in [r]} Z[j][h_j(i)] \leftarrow Z[j][h_j(i)] + s_j(i) \cdot \Delta$
- Query(i): output $x_i' = median_j Z[j][h_j(i)]$

For a fixed j:

$$\mathbb{E}[(x_i - Z[j][h_j(i)])^2] = \mathbb{E}\left[\left(\sum_{k \neq i} \Pr[h_j(k) = h_j(i)] x_k s_j(k)\right)^2\right] = \sum_{k \neq i} \frac{1}{t} x_k^2 \leq \frac{1}{t}|x|_2^2$$

$$\Pr\left[(x_i - Z[j][h_j(i)])^2 > \frac{3}{t}|x|_2^2\right] < \frac{1}{3}$$

$$\Pr\left[\left|x_i - Z[j][h_j(i)]\right| > \sqrt{\frac{3}{t}}|x|_2\right] < \frac{1}{3}$$

For $t = O\left(\frac{1}{\varepsilon^2}\right)$ we get

$$\Pr\left[\left|x_i - Z[j][h_j(i)]\right| > \varepsilon|x|_2\right] < \frac{1}{3}$$

so we get an L_2 norm guarantee. We then use the median with $r = O(\log \delta^{-1})$ to get $1 - \delta$ correctness.